

DecisionAssure + Sentinel

Strategic Business Expansion Blueprint for Runtime Governance, Enterprise Positioning, Pricing, and API-Based Collaboration

Purpose

This document proposes a practical commercial path in which DecisionAssure remains the execution-level trace and admissibility layer, while Sentinel becomes an optional external premium governance layer used when clients require deeper autonomous-system oversight. The goal is not to merge products, share code, or transfer intellectual property. The goal is to create a clean, independent, API-based value stack that can increase enterprise trust, pricing power, and market reach.

Core idea: DecisionAssure can open the door through execution traceability. Sentinel can raise the offer from a technical trace report to a higher-value runtime governance assessment.

1. Current Position vs. Combined Premium Position

Dimension	DecisionAssure Standalone	DecisionAssure + Sentinel
Client promise	Prove whether a specific execution was still admissible at commit time.	Prove execution integrity and assess whether the autonomous system is drifting, escalating, or becoming unstable over time.
Core value	Replayable trace, authority drift detection, rollback collapse, admissibility evidence.	Trace evidence plus trajectory governance, long-horizon drift, escalation analysis, orchestration risk, and governance recommendations.
Perceived category	Technical audit / runtime traceability.	Enterprise runtime governance and persistent autonomous oversight.
Best buyer	Technical founders, internal AI teams, compliance-aware startups.	Enterprise AI teams, regulated industries, autonomy programs, boards, risk/compliance leaders.
Typical pricing power	USD 1.5k - 3k for a focused pilot or trace report.	USD 5k - 15k for a combined assessment; USD 20k - 50k+ for larger enterprise reviews.

The commercial upgrade is simple: the client is no longer buying only a trace of what happened. The client is buying a governance view of what the system is becoming.

2. What Each Layer Should Own

Layer	Owner	Role	Must remain protected
DecisionAssure	Akhilesh Warik	Execution-bound admissibility, replayable traces, local authority drift, rollback viability, commit-time evidence.	DecisionAssure codebase, trace engine, methodology, client relationships.
Sentinel	Emanuele Colombo	External governance intelligence: trajectory risk, long-horizon drift, escalation modeling, orchestration oversight, governance scoring, recommendations.	Sentinel code, models, governance logic, GAI architecture, scoring methods, internal methodology.

Recommended Interoperability Model

- DecisionAssure remains fully independent and continues selling its core trace/admissibility work.
- Sentinel remains fully independent and is called only when an advanced governance layer is needed.
- Data exchange occurs through controlled API calls only: trace plus metadata in, governance analysis out.
- No code sharing, no internal integration, no transfer of methodology, and no dependency of Sentinel on DecisionAssure.
- DecisionAssure can be considered one compatible trace source for Sentinel, not the foundation of Sentinel.

Suggested API flow

Step	Action	Business meaning
1	Client requests a DecisionAssure assessment.	DecisionAssure owns the initial opportunity.
2	DecisionAssure generates trace and execution evidence.	The client receives local execution integrity.
3	If deeper governance is needed, DecisionAssure calls Sentinel API.	Sentinel activates as premium governance layer only when required.
4	Sentinel returns trajectory governance analysis and recommendations.	The offer becomes a higher-value governance assessment.
5	Client receives combined report.	Both parties monetize only the layer they contributed.

3. Pricing and Packaging

The key is to avoid pricing the combined stack like a small technical pilot. Once Sentinel is involved, the offer should be positioned as governance intelligence, not only a software trace.

Package	Scope	Target buyer	Suggested price	Revenue model
Package A - DecisionAssure Trace Pilot	One workflow, one trace, execution admissibility, short summary.	Startup AI teams, technical founders, internal AI teams.	USD 1.5k - 3k	100% DecisionAssure if Sentinel is not used.
Package B - Combined Runtime Governance Assessment	DecisionAssure trace plus Sentinel trajectory analysis, drift risk, escalation notes, governance score, recommendations.	AI agent startups, orchestration platforms, robotics startups, regulated SaaS.	USD 5k - 15k	70% Sentinel / 30% DecisionAssure when Sentinel contributes.
Package C - Enterprise Governance Review	Multiple workflows or agents, deeper governance map, long-horizon risks, orchestration risks, board-level recommendations.	Fintech, healthcare AI, critical operations, industrial automation, defense-adjacent autonomy.	USD 20k - 50k+	Engagement-specific split based on scope; default 70/30 if Sentinel is core layer.
Package D - Strategic Oversight Program	Recurring governance reviews, monthly trace analysis, Sentinel reports, improvement roadmap.	Enterprises running persistent autonomous systems.	USD 5k - 25k/month	Recurring split only when Sentinel is actively used.

Pricing logic

- A trace proves execution integrity. A governance assessment reduces enterprise risk.
- The combined offer should be priced higher because it addresses both what happened and where the system is heading.
- Sentinel should be charged as a premium governance layer, not as an add-on feature.
- Revenue share should apply only when Sentinel is actually used and contributes to the client deliverable.

4. Best Customer Targets

The best targets are not generic AI users. The best targets are teams building systems where autonomous decisions create operational, financial, legal, safety, or trust risk.

Priority	Customer type	Why they care	Best offer
1	AI agent orchestration startups	They need trust, governance, and differentiation to sell to enterprises.	Package B
2	Autonomous workflow platforms	They must prove decisions remain safe across multi-step execution.	Package B or C
3	Fintech and automated decision systems	They face audit, compliance, and risk requirements.	Package C
4	Healthcare AI governance teams	They need explainability, oversight, and safety around critical decisions.	Package C
5	Robotics and autonomous operations	Physical autonomy requires both execution safety and trajectory supervision.	Package C or D
6	Defense / critical infrastructure vendors	High-risk systems need persistent oversight and governance architecture.	Package C or D
7	Enterprise AI transformation teams	They need governance language and board-ready reports.	Package B or C

Recommended outreach angle

The strongest message is not: "we audit AI." The stronger message is: "we help prove whether autonomous workflows remain governable from individual execution to long-horizon behavior."

5. Commercial and Legal Guardrails

- Both systems remain sovereign and independent.
- No code sharing, no IP transfer, no internal integration.
- Each engagement should have a short written scope: client, workflow, deliverables, price, split, confidentiality, and liability cap.
- Default split for combined engagements: 70% Sentinel / 30% DecisionAssure when Sentinel is the premium governance layer.
- DecisionAssure standalone remains 100% DecisionAssure. Sentinel standalone remains 100% Sentinel.
- If either party only refers a client and does not contribute operational work, a simple 15% first-invoice referral model can be used instead.
- Sentinel usage should be tracked by client_id, engagement_id, API calls, reports generated, and agreed price.

Suggested next 30 days

Week	Action	Output
1	Finalize Sentinel trace intake API and DecisionAssure-compatible adapter.	Callable controlled endpoint for test traces.
1-2	Run one private combined demo using a simple multi-step workflow.	Internal proof of combined value.
2	Create a polished sample combined report.	Sales asset for future clients.
3	Identify 20 target companies in AI agents, orchestration, fintech, healthcare, robotics.	Prospect list.
4	Offer 3 paid combined assessments at early-adopter pricing.	First revenue and case studies.

Final positioning

DecisionAssure can be the execution evidence layer. Sentinel can be the premium governance intelligence layer. Together, they can offer a more complete runtime governance stack while remaining independent, protected, and commercially scalable.

This blueprint should be treated as a practical strategic suggestion for discussion, refinement, and controlled testing - not as a binding agreement or legal document.